


# Operating Systems

Key functions and structures of  
operating Systems.

CHAPTER 10



# Chapter Goals

- ▶ Describe the two main **responsibilities** of an operating system
- ▶ Define **memory** and **process** management
- ▶ Explain how **timesharing** creates the virtual machine illusion
- ▶ Explain the relationship between **logical** and **physical** addresses
- ▶ Compare and contrast **memory management techniques**

# Chapter Goals

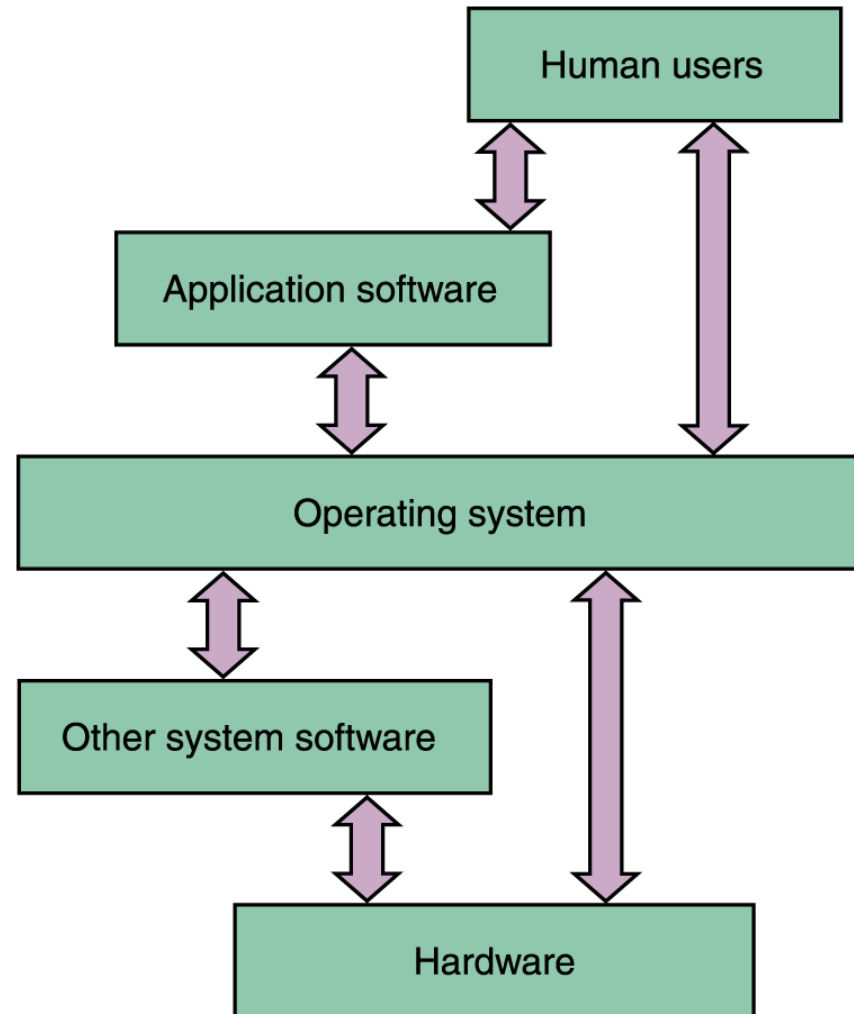
- ▶ Distinguish between **fixed** and **dynamic partitions**
- ▶ Define and apply partition **selection** algorithms
- ▶ Explain how **demand paging** creates the virtual memory illusion
- ▶ Explain the stages and transitions of the **process life cycle**
- ▶ Explain the processing of various CPU **scheduling** algorithms

# Roles of an Operating System

## Operating system

System software that

- **manages** computer resources, such as memory and input/output devices
- **provides** an interface through which a human can interact with the computer
- **allows** an application program to interact with these other system resources



# Roles of an Operating System

The various roles of an operating system generally revolve around the idea of “**sharing nicely**”

An operating system manages resources, and these resources are often shared in one way or another among programs that want to use them

# Resource Management

## Multiprogramming

The technique of keeping multiple programs that compete for access to the CPU in main memory at the same time so that they can execute

## Memory management

The process of keeping track of what programs are in memory and where in memory they reside

# Resource Management

## Process

A program in execution

## Process management

The act of carefully tracking the progress of a process and all of its intermediate states


## CPU scheduling

Determining which process in memory is executed by the CPU at any given point



# Batch Processing

The first operating system was a **human operator**, who organized various jobs from multiple users into *batches* of jobs that needed the same resources



# Timesharing

## Timesharing system

A system that allows multiple users to interact with a computer at the same time

## Virtual machine

The illusion created by a time-sharing system that each user has his/her own machine

# Other Factors

## Real-time System

A system in which response time is crucial given the nature of the application

## Response time

The time delay between receiving a stimulus and producing a response

## Device driver

A small program that “knows” the way a particular device expects to receive and deliver information

# Memory Management

Operating systems must employ techniques to

- Track where and how a program resides in memory
- Convert **logical addresses** into actual addresses

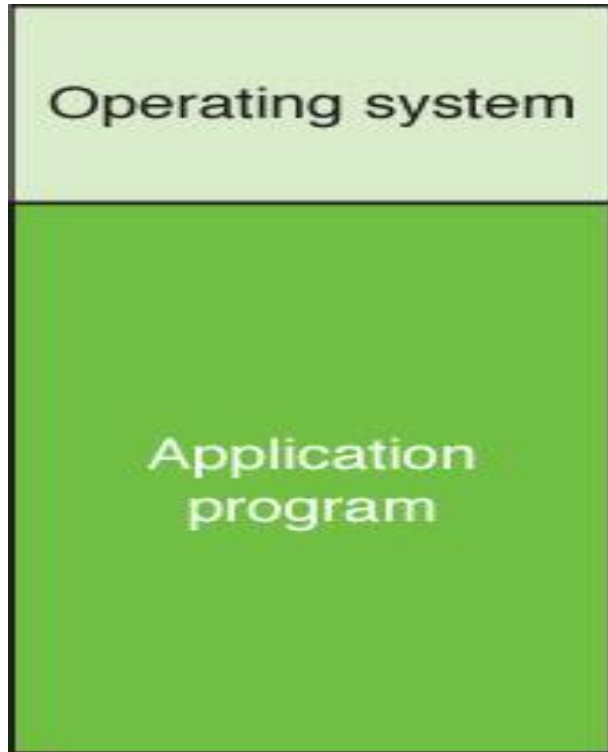
## Logical address

Reference to a stored value relative to the program making the reference

## Physical address

Actual address in main memory

# Single Contiguous MM



There are only two programs in memory

The operating system

The application program

This approach is called **single contiguous memory management**

**Figure 10.4** Main memory divided into two sections

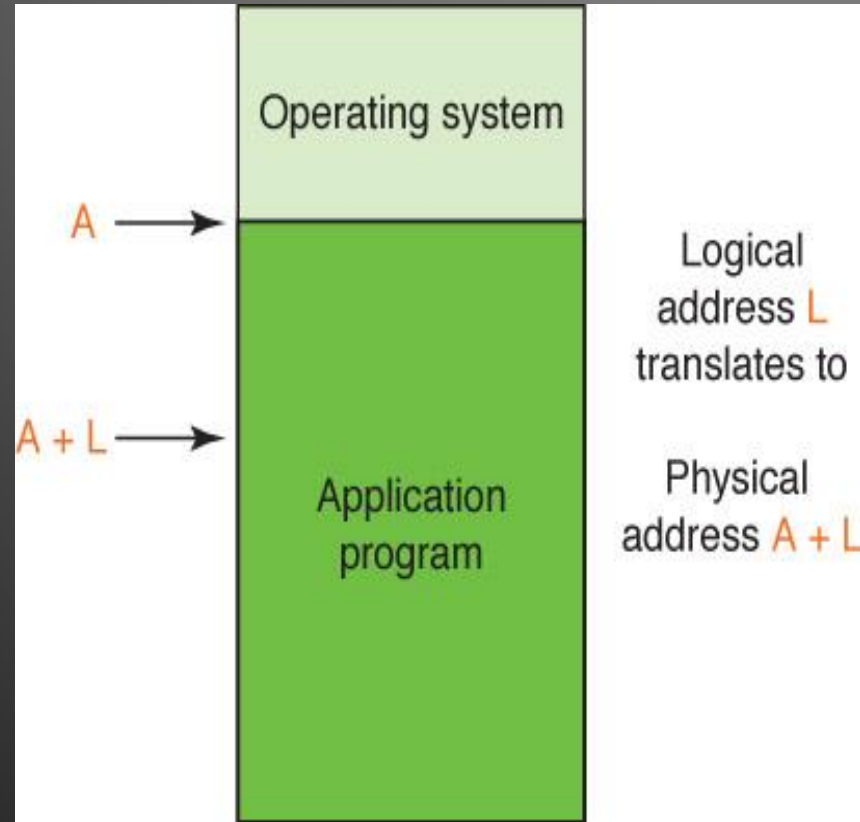
# Single Contiguous MM

## In concrete terms:

A **logical address** is simply an **integer** value relative to the **starting point** of the program

A **physical address** is a logical address added to the starting location of the program in main memory

# Single Contiguous MM



**Figure 10.5** Binding a logical address to a physical one

# Partition Memory Management

**Single contiguous MM** has only the OS and one other program in memory at one time

**Partition MM** has the OS and any number of other programs in memory at one time

There are two schemes for dividing up memory for programs:

- **Fixed partitions** Main memory is divided into a fixed number of partitions into which programs can be loaded
- **Dymanic partitions** Partitions are created as needed to fit the programs waiting to be loaded



# Partition Memory Management

Memory is divided into a set of partitions, some empty and some allocated to programs

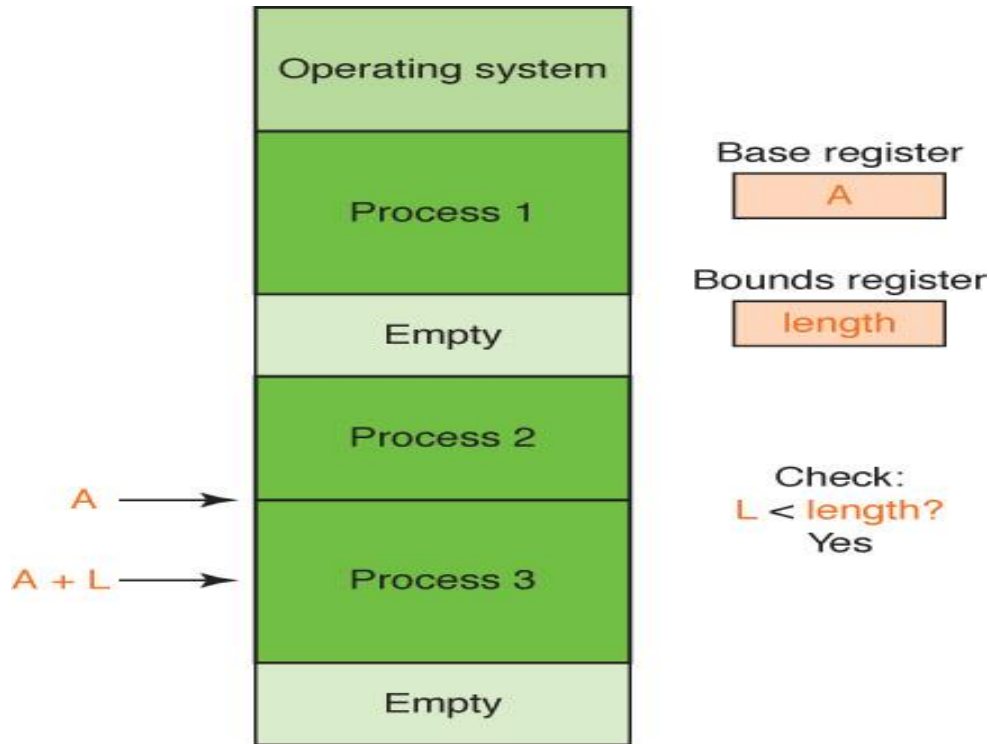
## Base register

A register that holds the beginning address of the current partition (the one that is running)

## Bounds register

A register that holds the length of the current partition

# Partition Memory Management



**Figure 10.6**  
Address resolution  
in partition memory  
management

# Partition Selection Algorithms

*Which partition should we allocate to a new program?*

- ▶ **First fit** Allocate program to the first partition big enough to hold it
- ▶ **Best fit** Allocated program to the smallest partition big enough to hold it
- ▶ **Worst fit** Allocate program to the largest partition big enough to hold it

# Partition Selection Algorithms

A: 1000
B: 700
C: 750
D: 1500
E: 300
F: 350

Requests come in for blocks of the following sizes:

1000, 25, 780, 1600, and 325

*What block will be assigned to each request if the*

- first-fit algorithm is used?*
- best-fit algorithm is used?*
- worst-fit algorithm is used?*

*(Treat each request as an independent event)*

# Paged Memory Management

## Paged memory technique

A technique in which processes are divided into fixed-size **pages** and stored in memory **frames** when loaded

### Frame

A fixed-size portion of *main memory* that holds a process page

### Page

A fixed-size portion of a *process* that is stored into a memory frame

**We assume that a frame and a page are the same size**

# Paged Memory Management

P1 PMT	
Page	Frame
0	5
1	12
2	15
3	7
4	22

P2 PMT	
Page	Frame
0	10
1	18
2	1
3	11

Memory	
Frame	Contents
0	
1	P2/Page2
2	
3	
4	
5	P1/Page0
6	
7	P1/Page3
8	
9	
10	P2/Page0
11	P2/Page3
12	P1/Page1
13	
14	
15	P1/Page2

—

### Figure 10.7

#### A paged memory management approach

Prog. 2, Page 2

Prog. 1, Page 3

# Paged Memory Management

P1 PMT	
Page	Frame
0	5
1	12
2	15
3	7
4	22

P2 PMT	
Page	Frame
0	10
1	18
2	1
3	11

Memory	
Frame	Contents
0	
1	P2/Page2
2	
3	
4	
5	P1/Page0
6	
7	P1/Page3
8	
9	
10	P2/Page0
11	P2/Page3
12	P1/Page1
13	
14	
15	P1/Page2

⋮

This new logical address is mapped to a physical address with the help of a page-map table (PMT)

Every program has a PMT that shows into which frame each page of the program is stored

*What is the physical address of <2, 518>?*

# Paged Memory Management

## Demand paging

An extension of paged memory management in which pages are brought into memory on demand

## Page swap

The act of bringing in a page from secondary memory, which often causes another page to be written back to secondary memory



# Paged Memory Management

## Virtual memory

The illusion that there are no restrictions on the size of a program because an entire process doesn't have to be in memory at the same time

## Thrashing

Inefficient processing caused by constant page swaps

# Process Management

## Process management

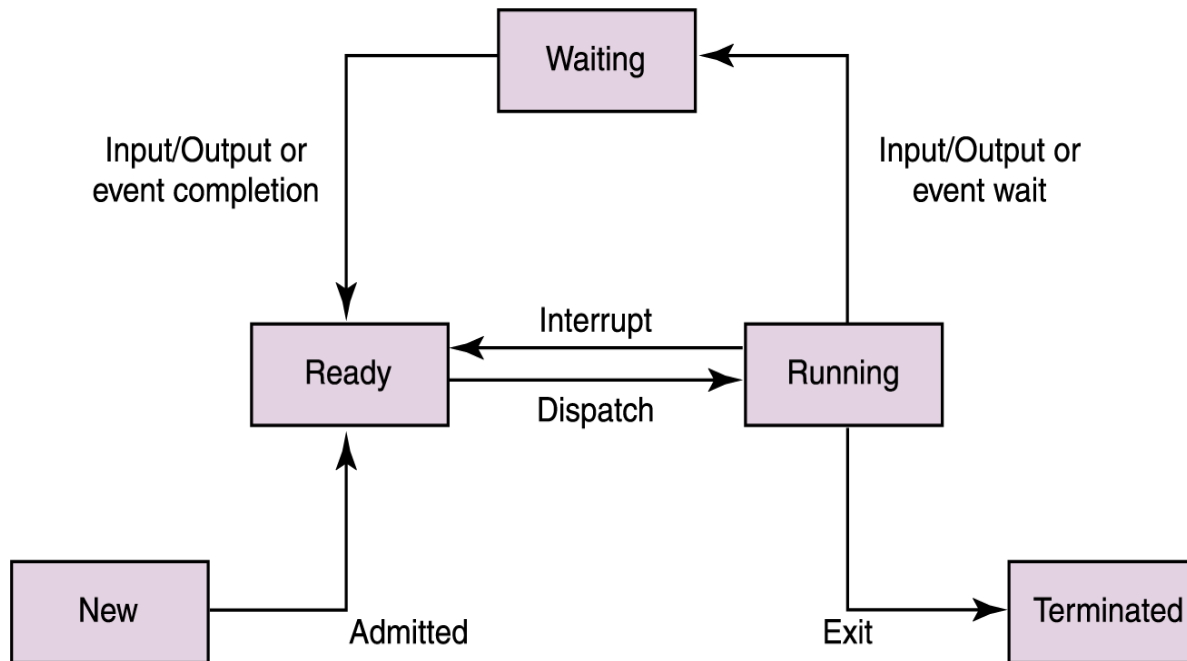
The act of managing the use of the CPU by individual processes

Recall that a process is a program in execution

*What stages does a process go through?*

# Process Management

## The Process States



**Figure 10.8** The process life cycle

# Process Management

## Process control block (PCB)

A *data structure* used by the OS to manage information about a process, including

- current value of the program counter
- values of all CPU registers for the process
- base and bound register values (or page tables)
- accounting information

Each *state* is represented by a list of PCBs, one for each process in that state

# Process Management

There is only one CPU and therefore only one set of CPU registers, which contain the values for the currently executing process

Each time a process is moved to the running state:

- Register values for the currently running process are stored into its PCB
- Its PCB is moved to the list of the state into which it goes
- Register values of the new process moving into the running state are loaded into the CPU
- This exchange of register information is called a **context switch**

# CPU Scheduling

## CPU Scheduling

The act of determining which process in the *ready* state should be moved to the *running* state

- Many processes may be in the ready state
- Only one process can be in the running state, making progress at any one time

*Which one gets to move from ready to running?*

# CPU Scheduling

## Nonpreemptive scheduling

The currently executing process gives up the CPU voluntarily

## Preemptive scheduling

The operating system decides to favor another process, preempting the currently executing process

## Turnaround time

The amount of time between when a process arrives in the ready state the first time and when it exits the running state for the last time

# CPU Scheduling Algorithms

## First-Come, First-Served

Processes are moved to the CPU in the order in which they arrive in the running state

## Shortest Job Next

Process with shortest estimated running time in the ready state is moved into the running state first


## Round Robin

Each process runs for a specified time slice and moves from the running state to the ready state to await its next turn if not finished



# First-Come, First-Served


Process	Service time
p1	140
p2	75
p3	320
p4	280
p5	125



Average turnaround time =  $(140+215+535+815+940)/5$   
=529

# Shortest Job Next (SJN)

Process	Service time
p1	140
p2	75
p3	320
p4	280
p5	125



$$\text{Average turnaround time} = (75 + 200 + 340 + 620 + 940) / 5 = 435$$

# Round Robin

Every process is treated the same!


## Time slice (quantum)

The amount of time each process receives before being preempted and returned to the ready state to allow another process its turn

# Round Robin

Suppose the time slice is 50

Process	Service time
p1	140
p2	75
p3	320
p4	280
p5	125



$$\begin{aligned}\text{Average turnaround time} &= (515 + 325 + 940 + 920 + 640) / 5 \\ &= 668\end{aligned}$$